



WEEK 16 — Mini Project + LLM Mindset (End-to-End NLP App)

🎯 Goal of Week 16 (In Simple Words)

Take the fine-tuned model from Week 15 and turn it into a **real working NLP application** with a simple UI (Gradio), plus evaluation + README.

✓ At the end, you will have a full “portfolio” project:

- Model
 - Inference pipeline
 - UI
 - Testing
 - Error analysis
 - Report / README
-

✅ STEP 0 — Choose Your Project (Pick One)

🅰️ Option A: Sentiment Analysis Chatbot ✅ (Recommended)

Input:

User types message/review

Output:

- Sentiment label (Positive/Negative/Uncertain)
- Confidence score
- Explanation (simple rule-based explanation)

✓ Best because:

- IMDB dataset is easy
- Labels are clean
- Works well with fine-tuned DistilBERT
- Great for beginners & submissions

▀ Option B: Fake News Detector (Advanced)

Input:

News headline or article

Output:

- Fake / Real probability

Harder because:

- Fake news datasets can be noisy
- “Real vs fake” can be subjective
- Need more cleaning + bias discussion

👉 I will write the steps in a way that works for both, but I'll assume **Option A** for example code (you can swap dataset/model later).

✓ STEP 1 — Project Folder Setup (Very Professional)

Create a clean project folder like this:

```
week16_nlp_project/
    |
    └── app/
        ├── app.py          # Gradio UI
        └── inference.py    # prediction function
    |
    └── model/
        └── fine_tuned_model/ # saved model + tokenizer (from week 15)
    |
    └── data/
        ├── sample_inputs.txt # real testing examples
        └── optional_dataset.csv # if you use custom data
    |
    └── results/
        └── predictions.csv
```

```
├── confusion_matrix.png  
└── error_analysis.md  
  
├── notebooks/  
│   ├── evaluation.ipynb  
│   └── testing.ipynb  
  
└── README.md  
└── requirements.txt
```

✓ Why this structure?

- Looks like real ML project
 - Easy to explain in interview
 - Easy to maintain
-

✓ STEP 2 — Load Fine-Tuned Model (From Week 15)

Definition (Beginner)

Inference means:

Using a trained model to make predictions (no training now).

Code: Load model + tokenizer

```
from transformers import AutoTokenizer, AutoModelForSequenceClassification  
import torch
```

```
MODEL_DIR = "./model/fine_tuned_model"
```

```
tokenizer = AutoTokenizer.from_pretrained(MODEL_DIR)  
model = AutoModelForSequenceClassification.from_pretrained(MODEL_DIR)  
model.eval()
```

Explanation (Line by line)

- `from_pretrained(MODEL_DIR)` loads your saved files
 - `model.eval()` makes predictions stable (no dropout)
-

STEP 3 — Build Inference Function (Core of the App)

What should inference function do?

Input text → output:

- label
 - confidence
 - probability scores (optional)
 - explanation
-

Code (Very Clean)

```
import torch

id2label = model.config.id2label

def predict(text):
    inputs = tokenizer(text, return_tensors="pt", truncation=True, padding=True)

    with torch.no_grad():
        outputs = model(**inputs)

    logits = outputs.logits
    probs = torch.softmax(logits, dim=1)[0]

    pred_id = torch.argmax(probs).item()
    label = id2label[pred_id]
    confidence = probs[pred_id].item()

    return label, confidence, probs.tolist()
```

Step-by-step explanation

- Tokenizer makes tensors
 - Model outputs logits
 - Softmax converts logits → probabilities
 - Argmax chooses best label
 - Confidence is probability of chosen label
-

STEP 4 — Add “Uncertain” Feature (Real ML Product Behavior)

Definition

A **threshold** is:

minimum confidence needed to accept prediction.

If confidence is low → show “Uncertain”.

Code

```
def predict_with_threshold(text, threshold=0.80):
    label, confidence, probs = predict(text)

    if confidence < threshold:
        return "UNCERTAIN", confidence, probs

    return label, confidence, probs
```

 This makes your app more realistic.

STEP 5 — Add Explanation (Simple but Very Important)

Transformers are black-box.

But for beginner projects, you can give:

- rule-based explanation
- highlight simple keywords

Simple explanation example (Sentiment)

- If positive words like “amazing, loved, great” → positive
 - If negative words like “worst, boring, hate” → negative
-

Code: Rule-based explanation

```
positive_words = ["love", "amazing", "great", "awesome", "fantastic", "best"]
negative_words = ["hate", "worst", "boring", "bad", "terrible", "waste"]

def simple_explanation(text):
    text_lower = text.lower()
    pos_hits = [w for w in positive_words if w in text_lower]
    neg_hits = [w for w in negative_words if w in text_lower]

    if pos_hits and not neg_hits:
        return f"Detected positive words: {pos_hits}"
    if neg_hits and not pos_hits:
        return f"Detected negative words: {neg_hits}"
    if pos_hits and neg_hits:
        return f"Mixed words found. Positive: {pos_hits}, Negative: {neg_hits}"
    return "No strong keywords detected. Model used full context to decide."
```

✓ In report, mention:

- this explanation is heuristic
 - not always perfect
-

✓ STEP 6 — Handle Long Text (Important for Fake News especially)

Transformers have max length (512 tokens).

Long text gets truncated.

Solution: Chunking

Split text → run prediction on each chunk → average.

Code: Chunking

```
def chunk_text(text, chunk_size=300):
    return [text[i:i+chunk_size] for i in range(0, len(text), chunk_size)]

def predict_long_text(text):
    chunks = chunk_text(text, 300)
```

```
all_probs = []
for c in chunks:
    _, _, probs = predict(c)
    all_probs.append(probs)

avg_probs = np.mean(all_probs, axis=0)
pred_id = int(np.argmax(avg_probs))
label = id2label[pred_id]
confidence = float(avg_probs[pred_id])

return label, confidence, avg_probs.tolist()
```

✓ Needed for:

- news articles
 - long reviews
 - paragraphs
-

✓ STEP 7 — Build Gradio UI (Very Simple & Powerful)

Definition

Gradio is a Python tool that creates UI for ML models quickly.

Install

pip install gradio

Gradio App Code (app.py)

import gradio as gr

```
def app_predict(text):
    label, confidence, probs = predict_with_threshold(text, threshold=0.80)
    explanation = simple_explanation(text)

    return {
        "Label": label,
        "Confidence": round(confidence, 4),
```

```

    "Probabilities": probs,
    "Explanation": explanation
}

demo = gr.Interface(
    fn=app_predict,
    inputs=gr.Textbox(lines=4, placeholder="Type your text here..."),
    outputs="json",
    title="Sentiment Analysis Chatbot (Fine-tuned DistilBERT)",
    description="Enter text → get sentiment + confidence + simple explanation."
)

demo.launch()

```

Explanation

- input = text box
 - output = JSON result
 - deploys locally in browser
-

STEP 8 — Testing with Real Examples (Must Do)

Create a file data/sample_inputs.txt and test:

- normal positive
- normal negative
- mixed sentiment
- sarcasm
- long text

Example test set

1. “This movie was fantastic and emotional.”
2. “Worst movie ever, total waste of time.”
3. “Good acting but terrible story.”
4. “Yeah great... I slept through it.” (sarcasm)

 Save outputs to results/predictions.csv.

STEP 9 — Error Analysis + Bias Discussion (Very Important)

What is Error Analysis?

Find wrong predictions and explain why.

Common reasons:

- sarcasm
- mixed sentiment
- long text
- domain mismatch

Bias Discussion (Simple)

- Model trained on movie reviews → may fail on product reviews
- Language style differences
- Non-English or Hinglish might be misclassified

 Write 8–10 lines about this in results/error_analysis.md.

STEP 10 — Write README (Final Submission)

README should include:

1. Project title
 2. What it does
 3. Model used
 4. Dataset used
 5. How to run (install + commands)
 6. Example inputs/outputs
 7. Limitations (bias + errors)
 8. Future improvements
-

STEP 11 — Final Checklist (Submission Ready

-  Fine-tuned model saved
 -  Inference function works
 -  Threshold logic added
 -  Explanation added
 -  Long text handling added
 -  Gradio UI running
 -  Testing examples done
 -  Error analysis written
 -  README written
 -  requirements.txt created
-

WEEK 16 DAILY PLAN (So You Know What To Do Each Day)

Day 1: Setup project folder + load model

Day 2: Build inference + threshold

Day 3: Add explanation logic

Day 4: Long text chunking

Day 5: Build Gradio UI

Day 6: Testing + CSV results

Day 7: Error analysis + README + final packaging
